

Parallel Programming in C with MPI and OpenMP

Michael J. Quinn



Chapter 11


Matrix Multiplication

Outline


- Sequential algorithms
 - ◆ Iterative, row-oriented
 - ◆ Recursive, block-oriented
- Parallel algorithms
 - ◆ Rowwise block striped decomposition
 - ◆ Cannon's algorithm

Iterative, Row-oriented Algorithm


Series of inner product (dot product) operations



Performance as n Increases




Reason: Matrix B Gets Too Big for Cache



Computing a row of C requires
accessing every element of B

Block Matrix Multiplication



Replace scalar multiplication
with matrix multiplication

Replace scalar addition with matrix addition

Recurse Until B Small Enough

A

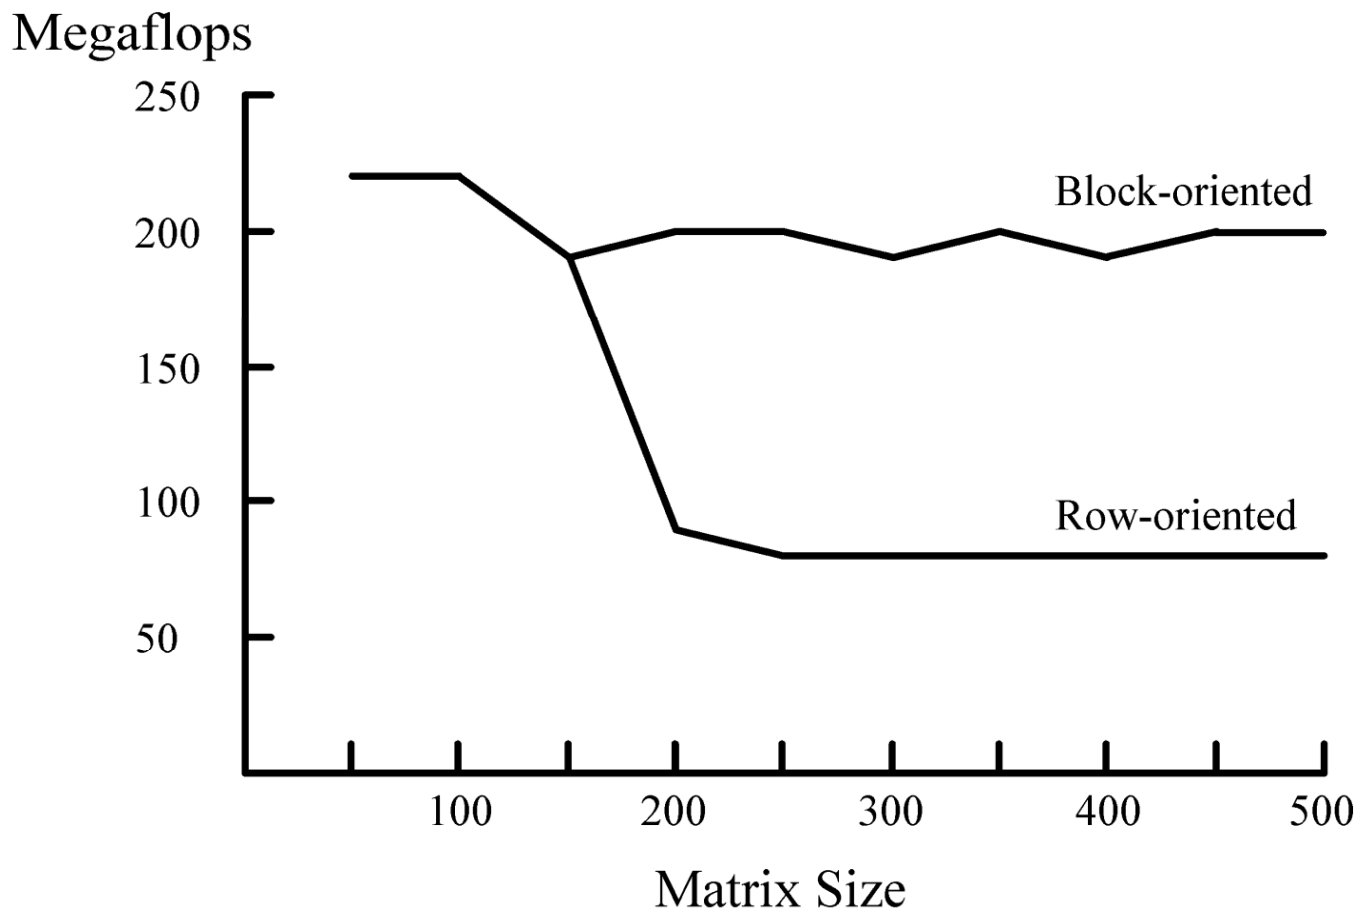
1	2	

B

1		
2		

C

Comparing Sequential Performance



First Parallel Algorithm

■ Partitioning

- ◆ Divide matrices into rows
- ◆ Each primitive task has corresponding rows of three matrices


■ Communication

- ◆ Each task must eventually see every row of B
- ◆ Organize tasks into a ring

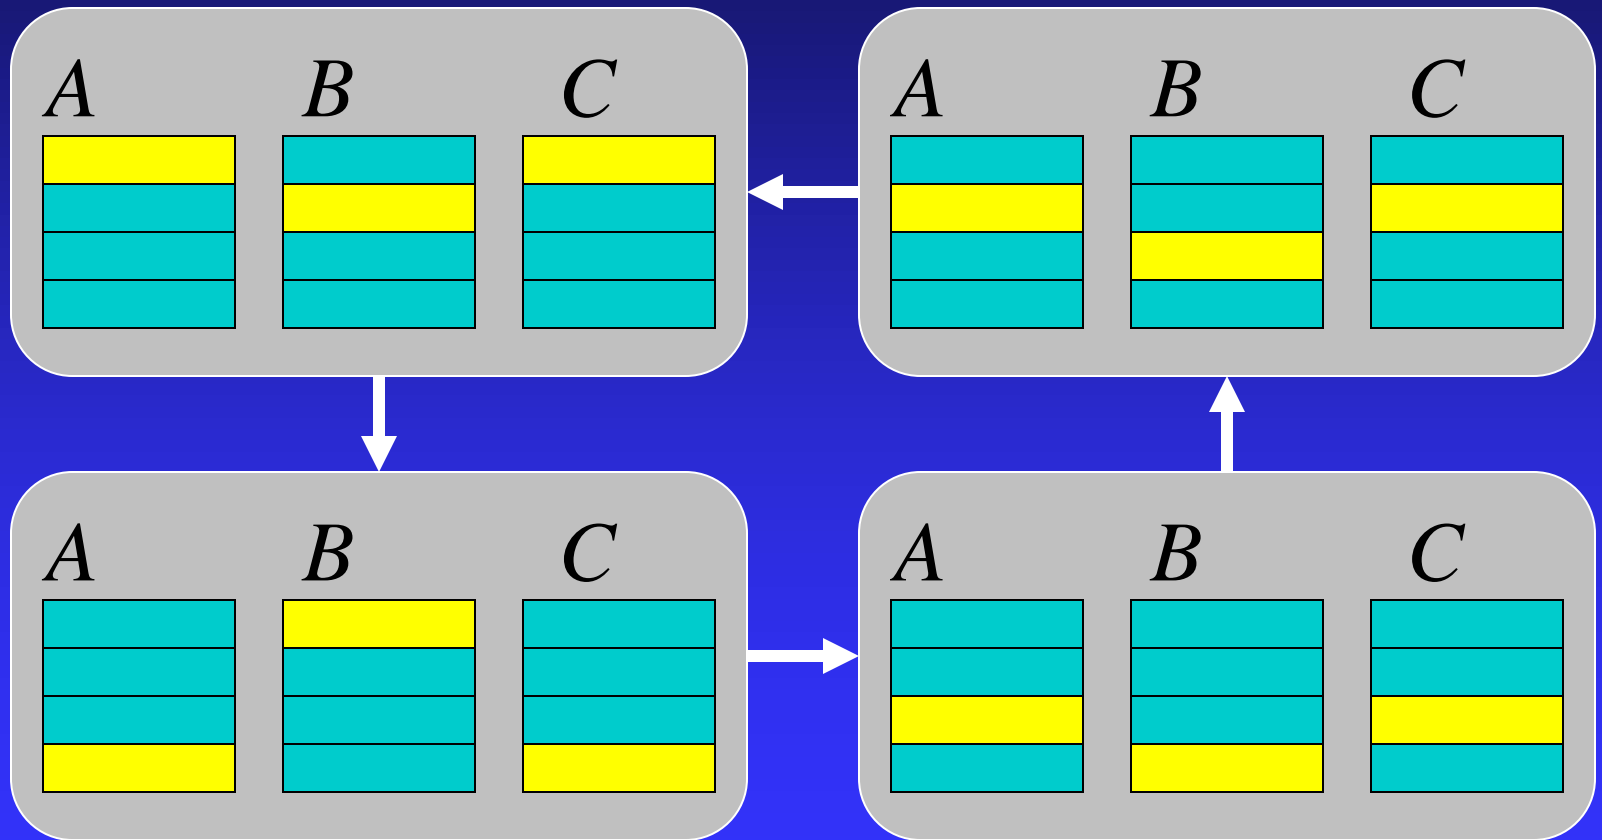
First Parallel Algorithm (cont.)

- Agglomeration and mapping
 - ◆ Fixed number of tasks, each requiring same amount of computation
 - ◆ Regular communication among tasks
 - ◆ Strategy: Assign each process a contiguous group of rows


Communication of B




Communication of B



Communication of B



Communication of B



Complexity Analysis

- Algorithm has p iterations
- During each iteration a process multiplies $(n / p) \times (n / p)$ block of A by $(n / p) \times n$ block of B: $\Theta(n^3 / p^2)$
- Total computation time: $\Theta(n^3 / p)$
- Each process ends up passing $(p-1)n^2/p = \Theta(n^2)$ elements of B

Isoefficiency Analysis

- Sequential algorithm: $\Theta(n^3)$
- Parallel overhead: $\Theta(pn^2)$

Isoefficiency relation: $n^3 \geq Cpn^2 \Rightarrow n \geq Cp$

$$M(Cp) / p = C^2 p^2 / p = C^2 p$$

- This system does not have good scalability


Weakness of Algorithm 1

- Blocks of B being manipulated have p times more columns than rows
- Each process must access every element of matrix B
- Ratio of computations per communication is poor: only $2n / p$


Parallel Algorithm 2 (Cannon's Algorithm)

- Associate a primitive task with each matrix element
- Agglomerate tasks responsible for a square (or nearly square) block of C
- Computation-to-communication ratio rises to n / \sqrt{p}

Elements of A and B Needed to Compute a Process's Portion of C



Algorithm 1



Cannon's
Algorithm

Blocks Must Be Aligned

$A_{0,0}$ $B_{0,0}$	$A_{0,1}$ $B_{0,1}$	$A_{0,2}$ $B_{0,2}$	$A_{0,3}$ $B_{0,3}$
------------------------	------------------------	------------------------	------------------------

$A_{1,0}$ $B_{1,0}$	$A_{1,1}$ $B_{1,1}$	$A_{1,2}$ $B_{1,2}$	$A_{1,3}$ $B_{1,3}$
------------------------	------------------------	------------------------	------------------------

$A_{2,0}$ $B_{2,0}$	$A_{2,1}$ $B_{2,1}$	$A_{2,2}$ $B_{2,2}$	$A_{2,3}$ $B_{2,3}$
------------------------	------------------------	------------------------	------------------------

$A_{3,0}$ $B_{3,0}$	$A_{3,1}$ $B_{3,1}$	$A_{3,2}$ $B_{3,2}$	$A_{3,3}$ $B_{3,3}$
------------------------	------------------------	------------------------	------------------------

$A_{0,0}$ $B_{0,0}$	$A_{0,1}$ $B_{1,1}$	$A_{0,2}$ $B_{2,2}$	$A_{0,3}$ $B_{3,3}$
------------------------	------------------------	------------------------	------------------------

$A_{1,1}$ $B_{1,0}$	$A_{1,2}$ $B_{2,1}$	$A_{1,3}$ $B_{3,2}$	$A_{1,0}$ $B_{0,3}$
------------------------	------------------------	------------------------	------------------------

$A_{2,2}$ $B_{2,0}$	$A_{2,3}$ $B_{3,1}$	$A_{2,0}$ $B_{0,2}$	$A_{2,1}$ $B_{1,3}$
------------------------	------------------------	------------------------	------------------------

$A_{3,3}$ $B_{3,0}$	$A_{3,0}$ $B_{0,1}$	$A_{3,1}$ $B_{1,2}$	$A_{3,2}$ $B_{2,3}$
------------------------	------------------------	------------------------	------------------------


Before

After


Blocks Need to Be Aligned

Each triangle represents a matrix block

Only same-color triangles should be multiplied




Rearrange Blocks



Block A_{ij} cycles
left i positions


Block B_{ij} cycles
up j positions

Consider Process $P_{1,2}$




Step 1

Consider Process $P_{1,2}$




Step 2

Consider Process $P_{1,2}$



Step 3

Consider Process $P_{1,2}$



Step 4

Complexity Analysis

- Algorithm has \sqrt{p} iterations
- During each iteration process multiplies two $(n / \sqrt{p}) \times (n / \sqrt{p})$ matrices: $\Theta(n^3 / p^{3/2})$
- Computational complexity: $\Theta(n^3 / p)$
- During each iteration process sends and receives two blocks of size $(n / \sqrt{p}) \times (n / \sqrt{p})$
- Communication complexity: $\Theta(n^2 / \sqrt{p})$

Isoefficiency Analysis

- Sequential algorithm: $\Theta(n^3)$
- Parallel overhead: $\Theta(\sqrt{p}n^2)$

Isoefficiency relation: $n^3 \geq C \sqrt{p}n^2 \Rightarrow n \geq C \sqrt{p}$

$$M(C\sqrt{p}) / p = C^2 p / p = C^2$$

- This system is highly scalable

Summary

- Considered two sequential algorithms
 - ◆ Iterative, row-oriented algorithm
 - ◆ Recursive, block-oriented algorithm
 - ◆ Second has better cache hit rate as n increases
- Developed two parallel algorithms
 - ◆ First based on rowwise block striped decomposition
 - ◆ Second based on checkerboard block decomposition
 - ◆ Second algorithm is scalable, while first is not