

Parallel Programming in C with MPI and OpenMP

Michael J. Quinn



Chapter 14

Sorting

Outline

- Sorting problem
- Sequential quicksort
- Parallel quicksort
- Hyperquicksort
- Parallel sorting by regular sampling

Sorting Problem

- Permute: unordered sequence \Rightarrow ordered sequence
- Typically key (value being sorted) is part of record with additional values (satellite data)
- Most parallel sorts designed for theoretical parallel models: not practical
- Our focus: internal sorts based on comparison of keys

Sequential Quicksort

17	14	65	4	22	63	11
----	----	----	---	----	----	----

Unordered list of values

Sequential Quicksort

17	14	65	4	22	63	11
----	----	----	---	----	----	----

Choose pivot value

Sequential Quicksort

14	4	11	17	65	22	63
----	---	----	----	----	----	----

Low list
(≤ 17)

High list
(> 17)

Sequential Quicksort

4	11	14	17	65	22	63
---	----	----	----	----	----	----

Recursively
apply quicksort
to low list

Sequential Quicksort

4	11	14	17	22	63	65
---	----	----	----	----	----	----

Recursively
apply quicksort
to high list

Sequential Quicksort

4	11	14	17	22	63	65
---	----	----	----	----	----	----

Sorted list of values

Attributes of Sequential Quicksort

- Average-case time complexity: $\Theta(n \log n)$
- Worst-case time complexity: $\Theta(n^2)$
 - ◆ Occurs when low, high lists maximally unbalanced at every partitioning step
- Can make worst-case less probable by using sampling to choose pivot value
 - ◆ Example: “Median of 3” technique

Quicksort Good Starting Point for Parallel Algorithm

■ Speed

- ◆ Generally recognized as fastest sort in average case
- ◆ Preferable to base parallel algorithm on fastest sequential algorithm

■ Natural concurrency

- ◆ Recursive sorts of low, high lists can be done in parallel

Definitions of “Sorted”

- Definition 1: Sorted list held in memory of a single processor
- Definition 2:
 - ◆ Portion of list in every processor's memory is sorted
 - ◆ Value of last element on P_i 's list is less than or equal to value of first element on P_{i+1} 's list
- We adopt Definition 2: Allows problem size to scale with number of processors

Parallel Quicksort

75, 91, 15, 64, 21, 8, 88, 54

P_0

50, 12, 47, 72, 65, 54, 66, 22

P_1

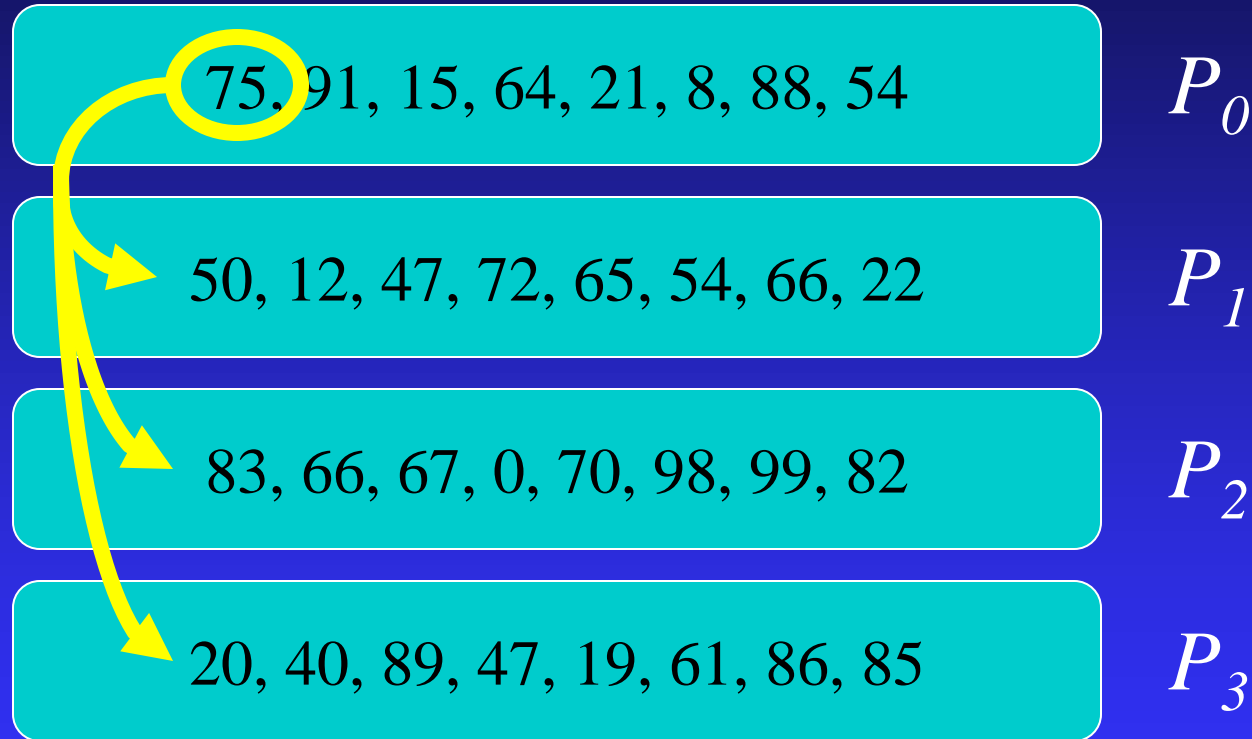
83, 66, 67, 0, 70, 98, 99, 82

P_2

20, 40, 89, 47, 19, 61, 86, 85


P_3

Parallel Quicksort




Process P_0 chooses and broadcasts randomly chosen pivot value

Parallel Quicksort




Exchange “lower half” and “upper half” values”

Parallel Quicksort




After exchange step

Parallel Quicksort



Processes P_0 and P_2 choose and broadcast randomly chosen pivots

Parallel Quicksort



Exchange values

Parallel Quicksort

Lower “half”
of lower “half”

15, 21, 8, 0, 12, 20, 19

P_0

Upper “half”
of lower “half”

50, 47, 72, 65, 54, 66, 22, 40,
47, 61, 75, 64, 54, 66, 67, 70

P_1

Lower “half”
of upper “half”

83, 82, 91, 88, 89, 86, 85

P_2

Upper “half”
of upper “half”

98, 99

P_3

Exchange values

Parallel Quicksort

Lower “half”
of lower “half”

0, 8, 12, 15, 19, 20, 21

P_0

Upper “half”
of lower “half”

22, 40, 47, 47, 50, 54, 54, 61,
64, 65, 66, 66, 67, 70, 72, 75

P_1

Lower “half”
of upper “half”

82, 83, 85, 86, 88, 89, 91

P_2

Upper “half”
of upper “half”

98, 99

P_3

Each processor sorts values it controls

Analysis of Parallel Quicksort

- Execution time dictated by when last process completes
- Algorithm likely to do a poor job balancing number of elements sorted by each process
- Cannot expect pivot value to be true median
- Can choose a better pivot value

Hyperquicksort

- Start where parallel quicksort ends: each process sorts its sublist
- First “sortedness” condition is met
- To meet second, processes must still exchange values
- Process can use median of its sorted list as the pivot value
- This is much more likely to be close to the true median

Hyperquicksort

75, 91, 15, 64, 21, 8, 88, 54

P_0

50, 12, 47, 72, 65, 54, 66, 22

P_1

83, 66, 67, 0, 70, 98, 99, 82

P_2

20, 40, 89, 47, 19, 61, 86, 85

P_3

Number of processors is a power of 2

Hyperquicksort

8, 15, 21, 54, 64, 75, 88, 91

P_0

12, 22, 47, 50, 54, 65, 66, 72

P_1

0, 66, 67, 70, 82, 83, 98, 99


P_2

19, 20, 40, 47, 61, 85, 86, 89

P_3


Each process sorts values it controls

Hyperquicksort



Process P_0 broadcasts its median value

Hyperquicksort



Processes will exchange “low”, “high” lists

Hyperquicksort

0, 8, 15, 21, 54

P_0

12, 19, 20, 22, 40, 47, 47, 50, 54

P_1

64, 66, 67, 70, 75, 82, 83, 88, 91, 98, 99


P_2

61, 65, 66, 72, 85, 86, 89

P_3


Processes merge kept and received values.

Hyperquicksort



Processes P_0 and P_2 broadcast median values.

Hyperquicksort



Communication pattern for second exchange

Hyperquicksort

0, 8, 12, 15

P_0

19, 20, 21, 22, 40, 47, 47, 50, 54, 54

P_1

61, 64, 65, 66, 66, 67, 70, 72, 75, 82

P_2

83, 85, 86, 88, 89, 91, 98, 99

P_3

After exchange-and-merge step

Complexity Analysis

Assumptions

- Average-case analysis
- Lists stay reasonably balanced
- Communication time dominated by message transmission time, rather than message latency

Complexity Analysis

- Initial quicksort step has time complexity $\Theta((n/p) \log (n/p))$
- Total comparisons needed for $\log p$ merge steps: $\Theta((n/p) \log p)$
- Total communication time for $\log p$ exchange steps: $\Theta((n/p) \log p)$

Isoefficiency Analysis

- Sequential time complexity: $\Theta(n \log n)$
- Parallel overhead: $\Theta(n \log p)$
- Isoefficiency relation:
 $n \log n \geq C n \log p \Rightarrow \log n \geq C \log p \Rightarrow n \geq p^C$

$$M(p^C) / p = p^C / p = p^{C-1}$$

- The value of C determines the scalability.
Scalability depends on ratio of communication speed to computation speed.

Another Scalability Concern

- Our analysis assumes lists remain balanced
- As p increases, each processor's share of list decreases
- Hence as p increases, likelihood of lists becoming unbalanced increases
- Unbalanced lists lower efficiency
- Would be better to get sample values from all processes before choosing median

Parallel Sorting by Regular Sampling (PSRS Algorithm)

- Each process sorts its share of elements
- Each process selects regular sample of sorted list
- One process gathers and sorts samples, chooses pivot values from sorted sample list, and broadcasts these pivot values
- Each process partitions its list into p pieces, using pivot values
- Each process sends partitions to other processes
- Each process merges its partitions

PSRS Algorithm

75, 91, 15, 64, 21, 8, 88, 54

P_0

50, 12, 47, 72, 65, 54, 66, 22

P_1

83, 66, 67, 0, 70, 98, 99, 82

P_2

Number of processors does not
have to be a power of 2.

PSRS Algorithm

8, 15, 21, 54, 64, 75, 88, 91

P_0

12, 22, 47, 50, 54, 65, 66, 72

P_1

0, 66, 67, 70, 82, 83, 98, 99

P_2

Each process sorts its list using quicksort.

PSRS Algorithm

8, 15, 21, 54, 64, 75, 88, 91

P_0

12, 22, 47, 50, 54, 65, 66, 72

P_1

0, 66, 67, 70, 82, 83, 98, 99

P_2

Each process chooses p regular samples.

PSRS Algorithm

8, 15, 21, 54, 64, 75, 88, 91

P_0

12, 22, 47, 50, 54, 65, 66, 72

P_1

0, 66, 67, 70, 82, 83, 98, 99

P_2

15, 54, 75, 22, 50, 65, 66, 70, 83

One process collects p^2 regular samples.

PSRS Algorithm

8, 15, 21, 54, 64, 75, 88, 91

P_0

12, 22, 47, 50, 54, 65, 66, 72

P_1

0, 66, 67, 70, 82, 83, 98, 99

P_2

15, 22, 50, 54, 65, 66, 70, 75, 83

One process sorts p^2 regular samples.

PSRS Algorithm

8, 15, 21, 54, 64, 75, 88, 91

P_0

12, 22, 47, 50, 54, 65, 66, 72

P_1


0, 66, 67, 70, 82, 83, 98, 99

P_2

15, 22, 50, 54, 65, 66, 70, 75, 83

One process chooses $p-1$ pivot values.

PSRS Algorithm



One process broadcasts $p-1$ pivot values.

PSRS Algorithm

8, 15, 21, 54, 64, 75, 88, 91

P_0

12, 22, 47, 50, 54, 65, 66, 72


P_1

0, 66, 67, 70, 82, 83, 98, 99

P_2

Each process divides list, based on pivot values.

PSRS Algorithm



Each process sends partitions to correct destination process.

PSRS Algorithm

0, 8, 12, 15, 21, 22, 47, 50

P_0

54, 54, 64, 65, 66, 66

P_1

67, 70, 72, 75, 82, 83, 88, 91, 98, 99

P_2

Each process merges p partitions.

Assumptions

- Each process ends up merging close to n/p elements
- Experimental results show this is a valid assumption
- Processor interconnection network supports p simultaneous message transmissions at full speed
- 4-ary hypertree is an example of such a network

Time Complexity Analysis

■ Computations

- ◆ Initial quicksort: $\Theta((n/p)\log(n/p))$
- ◆ Sorting regular samples: $\Theta(p^2 \log p)$
- ◆ Merging sorted sublists: $\Theta((n/p)\log p)$
- ◆ Overall: $\Theta((n/p)(\log n + \log p) + p^2 \log p)$

■ Communications

- ◆ Gather samples, broadcast pivots: $\Theta(\log p)$
- ◆ All-to-all exchange: $\Theta(n/p)$
- ◆ Overall: $\Theta(n/p)$

Isoefficiency Analysis

- Sequential time complexity: $\Theta(n \log n)$
- Parallel overhead: $\Theta(n \log p)$
- Isoefficiency relation:
$$n \log n \geq Cn \log p \Rightarrow \log n \geq C \log p$$
- Scalability function same as for hyperquicksort
- Scalability depends on ratio of communication to computation speeds

Summary

- Three parallel algorithms based on quicksort
- Keeping list sizes balanced
 - ◆ Parallel quicksort: poor
 - ◆ Hyperquicksort: better
 - ◆ PSRS algorithm: excellent
- Average number of times each key moved:
 - ◆ Parallel quicksort and hyperquicksort: $\log p / 2$
 - ◆ PSRS algorithm: $(p-1)/p$